

Problem 1.

- (a) N is the number of unordered pairs of distinct keys in $[U]$. This means $N = \binom{U}{2}$. Next, we want to find the expected number of unordered pairs of distinct keys x and y such that $h(x) = h(y)$ when we pick a random hash function h from \mathcal{H} (Given that \mathcal{H} is c -universal).

From the definition of c -universal, we know that $\Pr_{h \in \mathcal{H}}[h(x) = h(y)] \leq \frac{c}{m}$.

Let I_n be an indicator random variable that is 1 when $h(x) = h(y)$ and 0 otherwise for any given unordered pair of distinct keys x and y .

From this, we can determine $\mathbb{E}[I_n] \leq \frac{c}{m}$.

Let I be the number of unordered pairs of distinct keys x and y where $h(x) = h(y)$.

We can now express $\mathbb{E}[I]$:

$$\begin{aligned} I &= \sum_{x \neq y} I_n \\ \mathbb{E}[I] &= \sum_{x \neq y} \mathbb{E}[I_n] \\ &\leq N \cdot \frac{c}{m} \\ &\leq \binom{U}{2} \cdot \frac{c}{m} \\ &\leq \frac{U \cdot (U-1) \cdot c}{2 \cdot m} \end{aligned}$$

- (b) Let's consider one hash function $h \in \mathcal{H}$. Consider that h maps s_1 number of distinct keys to value 1, and s_2 number of distinct keys to value 2, and so on, all the way to s_m . This means that for any value of s_i , there are exactly $\binom{s_i}{2}$ possible pairings of keys such that $h(x) = h(y)$.

This let's us express I differently:

$$\begin{aligned} I &= \sum_{i=1}^m \binom{s_i}{2} \\ &= \sum_{i=1}^m \frac{s_i \cdot (s_i - 1)}{2} \end{aligned}$$

- (c) Now let's figure out s_i . For any given s_i , we know that $\mathbb{E}[s_i] \geq \frac{U}{m}$. This is because we are trying to fit U keys into m spaces. We can now use the equation in part b to determine the following:


$$\begin{aligned}
I &= \sum_{i=1}^m \frac{s_i \cdot (s_i - 1)}{2} \\
\mathbb{E}[I] &= \sum_{i=1}^m \frac{\mathbb{E}[s_i] \cdot (\mathbb{E}[s_i] - 1)}{2} \\
&\geq m \cdot \frac{\frac{U}{m} \cdot (\frac{U}{m} - 1)}{2} \\
&\geq \frac{U \cdot (\frac{U}{m} - 1)}{2}
\end{aligned}$$

- (d) By combining the inequalities of part A and part C, we can define a bound for the expected number of unordered pairs of distinct keys x and y such that $h(x) = h(y)$ when we pick a random hash function $h \in \mathcal{H}$.

$$\frac{U \cdot (\frac{U}{m} - 1)}{2} \leq \mathbb{E}[I] \leq \frac{U \cdot (U - 1) \cdot c}{2 \cdot m}$$

Problem 2.

- (a) We want to show that $f(x) = e^{tx}$ is convex for $t > 0$. By definition, a convex function is also described such that all points along any line between two points on the function has a value greater than or equal to the function value underneath the points. $f(x) = e^{tx}$ is convex over all t by observation:



- (b) Let Z be a random variable with probability density function g in the interval $[0,1]$. $p = \mathbb{E}[Z]$. Let's also define a Bernoulli random variable such that $\Pr[X = 1] = p$. We want to show that for any convex function f , the following is true.

$$\mathbb{E}[f(Z)] \leq \mathbb{E}[f(X)]$$

Let's start with the left hand side (LHS). We know that $\mathbb{E}[f(Z)] = \int_0^1 f(t)g(t)dt$ and $\mathbb{E}[Z] = \int_0^1 tg(t)dt$. We can also express t as $(t)(1) + (1 - t)(0)$. With these in mind, we can start simplifying the LHS:

$$\begin{aligned}
LHS &= \mathbb{E}[f(Z)] \\
&= \int_0^1 f(t)g(t)dt \\
&= \int_0^1 f((t)(1) + (1-t)(0))g(t)dt
\end{aligned}$$

We can now use the definition of a convex function, which is that for $0 \leq \lambda \leq 1$, and convex function f , the following is true: $f(\lambda x + (1-\lambda)y) \leq \lambda f(x) + (1-\lambda)f(y)$. We can plug this into what we have so far:

$$\begin{aligned}
LHS &= \int_0^1 f((t)(1) + (1-t)(0))g(t)dt \\
&\leq \int_0^1 tf(1) + (1-t)f(0)g(t)dt \\
&\leq \int_0^1 tf(1)g(t)dt + \int_0^1 (1-t)f(0)g(t)dt \\
&\leq f(1) \int_0^1 tg(t)dt + f(0) \left[\int_0^1 g(t)dt - \int_0^1 tg(t)dt \right] \\
&\leq f(1)\mathbb{E}[Z] + f(0)[1 - \mathbb{E}[Z]] \\
&\leq f(1)p + f(0)(1-p)
\end{aligned}$$

Now let's start examining the right hand side (RHS). Because X is a Bernoulli random variable (aka discrete values), we can express $\mathbb{E}[f(X)]$ as a computed expression:

$$\begin{aligned}
RHS &= f(0)Pr[X=0] + f(1)Pr[X=1] \\
&= f(0)(1-p) + f(1)p
\end{aligned}$$

Combine the expressions we have and we get:

$$\begin{aligned}
LHS &\leq RHS \\
\mathbb{E}[f(Z)] &\leq \mathbb{E}[f(X)]
\end{aligned}$$

- (c) Let Y_1, \dots, Y_n be independent identical distributed random variables over $[0,1]$. Let $Y = \sum_i Y_i$. We want to show that for $\delta \leq 1$, $Pr[Y - E[Y] > \delta] \leq \exp(-\delta^2/2n)$.

We can use Hoeffding's equality here. In particular, if we let $\epsilon = \delta/n$, we get the above expression. The reasoning behind this is that we can treat each Y_i as an identical experiment. The variance of the experiment being conducted n times should decrease, but $E[Y]$ stays the same. Therefore, we can divide δ by n .

$$\begin{aligned}
Pr[X - E[X] \geq \epsilon n] &\leq \exp\left(\frac{-\epsilon^2 n}{2}\right) \\
\epsilon &= \delta/n \\
Pr[Y - E[Y] \geq (\delta/n)n] &\leq \exp\left(\frac{-(\delta/n)^2 n}{2}\right) \\
Pr[Y - E[Y] \geq \delta] &\leq \exp\left(\frac{-\delta^2}{2n}\right)
\end{aligned}$$

Problem 3.

- (a) We want to show that the hash functions in \mathcal{H} have the following property: for any key $x \in [U]$ and $v \in [m]$, we have

$$Pr_{h \in \mathcal{H}}[h(x) = v] = \frac{1}{m}$$

To show this, we can break down v in terms of the hash function definition. More specifically, we can relate the two together:

$$h(x) = H_0(x_0) \oplus H_1(x_1) \oplus \dots \oplus H_{c-1}(x_{c-1}) \quad (1)$$

$$v = v_0 \oplus v_1 \oplus \dots \oplus v_{c-1} \quad (2)$$

We can also express $Pr[h(x) = v]$ as the probability that each character is the correct hash:

$$Pr_{h \in \mathcal{H}}[h(x) = v] = \prod_{i=0}^{c-1} Pr[H_i[x_i] = v_i]$$

Individually, $Pr[H_i[x_i] = v_i]$ is $\frac{1}{m^{1/c}}$. We know this because the size of each H_i is $m^{1/c}$, and we are trying to select an individual entry. Overall, then, we can compute the above expression to arrive at our answer

$$\begin{aligned}
Pr_{h \in \mathcal{H}}[h(x) = v] &= \prod_{i=0}^{c-1} Pr[H_i[x_i] = v_i] \\
&= \left(\frac{1}{m^{1/c}}\right)^c \\
&= \frac{1}{m}
\end{aligned}$$

- (b) We want to show that for two different keys $x, y \in [U]$ and $u, v \in [m]$, the following property is true:

$$Pr_{h \in \mathcal{H}}[h(x) = u \text{ and } h(y) = v] = \frac{1}{m^2}$$

Intuitively, because x and y are distinct keys, we can regard the above expression as the product of two independent events.

$$\Pr_{h \in \mathcal{H}}[h(x) = u \text{ and } h(y) = v] = \Pr_{h \in \mathcal{H}}[h(x) = u] \times \Pr_{h \in \mathcal{H}}[h(y) = v]$$

The probability of each event happening independently together is the same as the probability we computed above:

$$\begin{aligned} \Pr_{h \in \mathcal{H}}[h(x) = u] &= \frac{1}{m} \\ \Pr_{h \in \mathcal{H}}[h(y) = v] &= \frac{1}{m} \\ \Pr_{h \in \mathcal{H}}[h(x) = u \text{ and } h(y) = v] &= \frac{1}{m} \times \frac{1}{m} = \frac{1}{m^2} \end{aligned}$$

(c) Intentionally left blank

(d) Suppose $c = 2$. Imagine we have four distinct keys w, x, y, z , and the hash values for three of them (any three) r, s, t . We want to determine the hash value of the last key u . This is possible if we closely examine the possible values for the keys and hashes:

$$\begin{aligned} h(w) &= H_0[w_0] \oplus H_1[w_1] \\ h(x) &= H_0[x_0] \oplus H_1[x_1] \\ h(y) &= H_0[y_0] \oplus H_1[y_1] \\ h(z) &= H_0[z_0] \oplus H_1[z_1] \end{aligned}$$

At first glance, the characters $w_0, w_1, x_0, x_1, y_0, y_1, z_0, z_1$ have nothing to do with each other. But, because $c = 2$, we can actually constraint the characters:

$$\begin{aligned} w_0, x_0, y_0, z_0 &\in \{c_0^{(0)}, c_1^{(0)}\} \\ w_1, x_1, y_1, z_1 &\in \{c_0^{(1)}, c_1^{(1)}\} \end{aligned}$$

Therefore:

$$(k_0, k_1) \in \{(c_0^{(0)}, c_0^{(1)}), (c_0^{(0)}, c_1^{(1)}), (c_1^{(0)}, c_0^{(1)}), (c_1^{(0)}, c_1^{(1)})\} \forall k \in \{w, x, y, z\}$$

With these constraints, we can also breakdown the hash values:

$$h_0 = H_0[c_0^{(0)}] \oplus H_1[c_0^{(1)}] \quad (3)$$

$$h_1 = H_0[c_0^{(0)}] \oplus H_1[c_1^{(1)}] \quad (4)$$

$$h_2 = H_0[c_0^{(1)}] \oplus H_1[c_0^{(1)}] \quad (5)$$

$$h_3 = H_0[c_0^{(1)}] \oplus H_1[c_1^{(1)}] \quad (6)$$

$$r, s, t, u \in \{h_0, h_1, h_2, h_3\}$$

With r, s, t, u being distinct, once we have three keys, all we need to do is find the missing pair of characters. This can be achieved with an interesting observation about h_0, h_1, h_2, h_3 , which is that they represent the set of all possible hashes. Therefore:

$$h_0 \oplus h_1 \oplus h_2 \oplus h_3 = 0 \quad (7)$$

$$r \oplus s \oplus t \oplus u = 0 \quad (8)$$

- (e) We want to show that for any set of d different keys $x^{(1)}, x^{(2)}, \dots, x^{(d)}$, there exists an index $i \in [c]$ such that the i th characters of those keys share at least $d^{1/c}$ different values.

We can show this via contradiction. The contradictory argument in this case is that $\forall i \in [c]$, the i th characters of those d keys share less than $d^{1/c}$ different values.

$$h(x^{(1)}) = H_1[x_1^{(1)}] \oplus H_2[x_2^{(1)}] \oplus \dots \oplus H_{c-1}[x_{c-1}^{(1)}] \quad (9)$$

$$h(x^{(2)}) = H_1[x_1^{(2)}] \oplus H_2[x_2^{(2)}] \oplus \dots \oplus H_{c-1}[x_{c-1}^{(2)}] \quad (10)$$

$$\vdots \quad (11)$$

$$h(x^{(d)}) = H_1[x_1^{(d)}] \oplus H_2[x_2^{(d)}] \oplus \dots \oplus H_{c-1}[x_{c-1}^{(d)}] \quad (12)$$

For any vertical slice $x_i^{(1)}, x_i^{(2)}, \dots, x_i^{(d)} \forall i \in [c]$, there are $< d^{1/c}$ distinct letters.

Let Y_i be the # of distinct letters for the i th character. From the contradictory argument above, they must share less than $d^{1/c}$ different values.

$$Y_i < d^{1/c}$$

Let Y be the # of distinct keys. From a tabulation hashing standpoint, a key is composed of its characters. In order for a key to be distinct, at least one of the characters must be distinct. This let's us express Y in terms of Y_i :

$$Y = \prod_{i=1}^{c-1} Y_i$$

Combining the two parts together yields:

$$\begin{aligned} Y &= \prod_{i=1}^{c-1} Y_i \\ &< (d^{1/c})^c \\ &< d \end{aligned}$$

But, from the original problem statement, the number of distinct keys is d . $Y = d$ and $Y < d$ creates a contradiction.

(f) Intentionally Left Blank

Total Time Taken: 20 Hours.